

In [5]: *#Importing the required libraries*

```
import keras
from keras.models import Sequential
from keras.layers import Convolution2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense
from keras.layers import Dropout
from keras.utils import np_utils

import matplotlib.pyplot as plt
```

In [6]: *#Load the CIFAR-10 data*

```
from keras.datasets import cifar10
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 155s 1us/step
170508288/170498071 [=====] - 155s 1us/step

In [7]: *#Display few images from the dataset*

```
for i in range(0, 9):
    plt.subplot(3,3,i + 1)
    plt.imshow(X_train[i])
# show the plot
plt.show()
```



In [8]: *#Convert the input into floats*

```
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
```

```
#Normalize the data such that the values are between 0 and 1
X_train = X_train / 255.0
X_test = X_test / 255.0
```

In [9]: *#one hot encode the outputs*

```
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
```

In [10]: *#Assingning values to hyperparameters*

```
lr_rate = 0.01
batch_size = 35
epochs = 20
n_classes = 10
```

In [11]: *#create the model*

```
model = Sequential()

#Layer 1
model.add(Convolution2D(32, kernel_size = (3,3), activation = 'relu', input_shape = (32, 32, 3)))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

#Layer 2
model.add(Convolution2D(64, kernel_size = (3,3), activation = 'relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

#Layer 3
model.add(Convolution2D(128, kernel_size = (3,3), activation = 'relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

model.add(Flatten()) #Flattens the previous layer into a huge vector

#Fully connected layer
model.add(Dense(512, activation = 'relu', kernel_initializer = 'uniform'))
model.add(Dropout(0.30))
model.add(Dense(10, activation = 'softmax'))
```

In [12]: *#compile the CNN*

```
#from keras.optimizers import SGD
#sgd = SGD(lr = lr_rate)
model.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics = ['accuracy'])
```

In [13]: *#Fit the CNN model to our dataset*

```
output = model.fit(X_train, y_train, batch_size = batch_size, epochs = epochs, validation_data = (X_test, y_test))
```

Epoch 1/20
1429/1429 [=====] - 83s 55ms/step - loss: 1.6143 - accuracy: 0.4050 - val_loss: 1.3111 - val_accuracy: 0.5351
Epoch 2/20
1429/1429 [=====] - 73s 51ms/step - loss: 1.2795 - accuracy: 0.5399 - val_loss: 1.0963 - val_accuracy: 0.6049
Epoch 3/20
1429/1429 [=====] - 78s 54ms/step - loss: 1.1521 - accuracy: 0.5902 - val_loss: 1.0256 - val_accuracy: 0.6318
Epoch 4/20
1429/1429 [=====] - 71s 50ms/step - loss: 1.0714 - accuracy: 0.6199 - val_loss: 0.9447 - val_accuracy: 0.6685
Epoch 5/20
1429/1429 [=====] - 71s 49ms/step - loss: 1.0134 - accuracy: 0.6437 - val_loss: 0.9189 - val_accuracy: 0.6815
Epoch 6/20
1429/1429 [=====] - 68s 47ms/step - loss: 0.9681 - accuracy: 0.6608 - val_loss: 0.9097 - val_accuracy: 0.6832
Epoch 7/20
1429/1429 [=====] - 65s 45ms/step - loss: 0.9377 - accuracy: 0.6712 - val_loss: 0.8387 - val_accuracy: 0.7116
Epoch 8/20
1429/1429 [=====] - 71s 50ms/step - loss: 0.9093 - accuracy: 0.6825 - val_loss: 0.8265 - val_accuracy: 0.7127
Epoch 9/20
1429/1429 [=====] - 75s 52ms/step - loss: 0.8835 - accuracy: 0.6878 - val_loss: 0.8011 - val_accuracy: 0.7183
Epoch 10/20
1429/1429 [=====] - 102s 72ms/step - loss: 0.8647 - accuracy: 0.6956 - val_loss: 0.8145 - val_accuracy: 0.7154
Epoch 11/20
1429/1429 [=====] - 65s 46ms/step - loss: 0.8530 - accuracy: 0.6996 - val_loss: 0.8080 - val_accuracy: 0.7209
Epoch 12/20
1429/1429 [=====] - 68s 47ms/step - loss: 0.8389 - accuracy: 0.7083 - val_loss: 0.7912 - val_accuracy: 0.7320
Epoch 13/20
1429/1429 [=====] - 67s 47ms/step - loss: 0.8212 - accuracy: 0.7126 - val_loss: 0.8288 - val_accuracy: 0.7104
Epoch 14/20
1429/1429 [=====] - 72s 50ms/step - loss: 0.8114 - accuracy: 0.7151 - val_loss: 0.7476 - val_accuracy: 0.7461
Epoch 15/20
1429/1429 [=====] - 78s 55ms/step - loss: 0.7994 - accuracy: 0.7185 - val_loss: 0.7674 - val_accuracy: 0.7331
Epoch 16/20
1429/1429 [=====] - 89s 62ms/step - loss: 0.7864 - accuracy: 0.7226 - val_loss: 0.7427 - val_accuracy: 0.7479
Epoch 17/20
1429/1429 [=====] - 130s 91ms/step - loss: 0.7828 - accuracy: 0.7262 - val_loss: 0.7699 - val_accuracy: 0.7407
Epoch 18/20
1429/1429 [=====] - 127s 89ms/step - loss: 0.7800 - accuracy: 0.7287 - val_loss: 0.7302 - val_accuracy: 0.7513
Epoch 19/20
1429/1429 [=====] - 126s 88ms/step - loss: 0.7649 - accuracy: 0.7310 - val_loss: 0.7485 - val_accuracy: 0.7405
Epoch 20/20
1429/1429 [=====] - 125s 87ms/step - loss: 0.7676 - accuracy: 0.7327 - val_loss: 0.7465 - val_accuracy: 0.7441

In [14]: print(model.evaluate(X_test, y_test))

313/313 [=====] - 7s 23ms/step - loss: 0.7465 - accuracy: 0.7441
[0.7464656829833984, 0.7440999746322632]

In [17]: *# Loss Curves*

```
plt.figure(figsize=[14,10])
plt.subplot(211)
plt.plot(output.history['loss'], 'r', linewidth=3.0)
plt.plot(output.history['val_loss'], 'b', linewidth=3.0)
plt.legend(['Training loss', 'Validation Loss'], fontsize=18)
plt.xlabel('Epochs ', fontsize=16)
plt.ylabel('Loss', fontsize=16)
plt.title('Loss Curves', fontsize=16)

# Accuracy Curves
plt.figure(figsize=[14,10])
plt.subplot(212)
plt.plot(output.history['accuracy'], 'r', linewidth=3.0)
plt.plot(output.history['val_accuracy'], 'b', linewidth=3.0)
plt.legend(['Training Accuracy', 'Validation Accuracy'], fontsize=18)
plt.xlabel('Epochs ', fontsize=16)
plt.ylabel('Accuracy', fontsize=16)
plt.title('Accuracy Curves', fontsize=16)
```

Out[17]: Text(0.5, 1.0, 'Accuracy Curves')

